

## Tanzanian Well Data Analysis Notebook

The following code will attempt to find correlations between the available data to help determine well functionality in the nation of Tanzania. Several methods are recommended to do this successfully, to most accurately achieve results. Since this is quite literally a life or death data problem, careful analysis should be used throughout the process.

```
In [405]: #Import the relevant libraries
import pandas as pd
from __future__ import print_function, division
import pandas as pd
import numpy as np
import os
from sklearn.metrics import accuracy_score
%matplotlib inline
import matplotlib.pyplot as plt
from sklearn.datasets import make_classification
from sklearn.metrics import accuracy_score, classification_report
from sklearn.linear_model import LogisticRegression
import warnings
from sklearn.exceptions import DataConversionWarning
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from scipy.cluster import hierarchy as hc
```

## OSEMN-Obtain

Here is where we will import the data into our notebook as csv files, generously provided to us by Taarifa, and the Tanzanian Ministry of Water

```
In [299]: #import the data as csv files
submission_format = pd.read_csv('SubmissionFormat.csv')
test = pd.read_csv('702ddfc5-68cd-4d1d-a0de-f5f566f76d91.csv')
labels = pd.read_csv('0bf8bc6e-30d0-4c50-956a-603fc693d966.csv')
train = pd.read_csv('4910797b-ee55-40a7-8668-10efd5c1b960.csv')
```

```
In [300]: #Check each data collection to make sure they're appropriately labeled
submission_format.head()
test.head()
train.head()
labels.head()
```

```
Out[300]:
```

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional

## OSEMN-Explore

Now that we have the data in our notebook, we can start executing functions to explore the size, shape, and types of data we will be working with.

```
In [301]: print("Train data labels:", len(labels))
print("Train data rows, columns:", train.shape)
print("Test data rows, columns:", test.shape)
```

```
Train data labels: 59400
Train data rows, columns: (59400, 40)
Test data rows, columns: (14850, 40)
```

```
In [302]: labels.status_group.value_counts()
```

```
Out[302]: functional          32259
non functional          22824
functional needs repair    4317
Name: status_group, dtype: int64
```

```
In [303]: labels.status_group.value_counts(normalize=True)
```


```
Out[303]: functional          0.543081
non functional          0.384242
functional needs repair    0.072677
Name: status_group, dtype: float64
```

```
In [304]: majority = labels['status_group'].mode()[0]
y_pred = np.full(shape=labels['status_group'].shape, fill_value=majority)
```

```
In [306]: all(y_pred==majority)
```

```
Out[306]: True
```

```
In [307]: plt.figure(figsize = (16,8))  
plt.title("Water Pumps Status",fontsize=20)  
sns.countplot(x = labels['status_group'], data = labels, palette="BuPu");
```



## OSEMN-Scrub

A quick look at the data shows us there are quite a few wells in need of repair, removal or relocation. Let's go a little further by removing missing values and changing some of our data types so we can have a clearer picture of what we'll be dealing with, and make a little less work for our model later on!

```
In [308]: df = train.merge(labels, on = train.index)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 43 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   key_0                                59400 non-null  int64
1   id_x                                 59400 non-null  int64
2   amount_tsh                          59400 non-null  float64
3   date_recorded                       59400 non-null  object
4   funder                              55765 non-null  object
5   gps_height                          59400 non-null  int64
6   installer                          55745 non-null  object
7   longitude                           59400 non-null  float64
8   latitude                           59400 non-null  float64
9   wpt_name                            59400 non-null  object
10  num_private                         59400 non-null  int64
11  basin                              59400 non-null  object
12  subvillage                         59029 non-null  object
13  region                             59400 non-null  object
14  region_code                        59400 non-null  int64
15  district_code                     59400 non-null  int64
16  lga                                59400 non-null  object
17  ward                               59400 non-null  object
18  population                         59400 non-null  int64
19  public_meeting                    56066 non-null  object
20  recorded_by                       59400 non-null  object
21  scheme_management                 55523 non-null  object
22  scheme_name                       31234 non-null  object
23  permit                           56344 non-null  object
24  construction_year                 59400 non-null  int64
25  extraction_type                   59400 non-null  object
26  extraction_type_group             59400 non-null  object
27  extraction_type_class             59400 non-null  object
28  management                       59400 non-null  object
29  management_group                 59400 non-null  object
30  payment                           59400 non-null  object
31  payment_type                     59400 non-null  object
32  water_quality                    59400 non-null  object
33  quality_group                    59400 non-null  object
34  quantity                         59400 non-null  object
35  quantity_group                   59400 non-null  object
36  source                           59400 non-null  object
37  source_type                      59400 non-null  object
38  source_class                     59400 non-null  object
39  waterpoint_type                  59400 non-null  object
40  waterpoint_type_group            59400 non-null  object
41  id_y                             59400 non-null  int64
42  status_group                     59400 non-null  object
dtypes: float64(3), int64(9), object(31)
memory usage: 19.9+ MB
```

```
In [309]: !pip install category_encoders
!pip install geopandas
import geopandas
```

```
Requirement already satisfied: category_encoders in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (2.2.2)
Requirement already satisfied: pandas>=0.21.1 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from category_encoders) (1.1.3)
Requirement already satisfied: statsmodels>=0.9.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from category_encoders) (0.12.0)
Requirement already satisfied: scipy>=1.0.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from category_encoders) (1.5.2)
Requirement already satisfied: patsy>=0.5.1 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from category_encoders) (0.5.1)
Requirement already satisfied: numpy>=1.14.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from category_encoders) (1.18.5)
Requirement already satisfied: scikit-learn>=0.20.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from category_encoders) (0.23.2)
Requirement already satisfied: python-dateutil>=2.7.3 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=0.21.1->category_encoders) (2.8.1)
Requirement already satisfied: pytz>=2017.2 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=0.21.1->category_encoders) (2020.1)
Requirement already satisfied: six in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from patsy>=0.5.1->category_encoders) (1.15.0)
Requirement already satisfied: joblib>=0.11 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from scikit-learn>=0.20.0->category_encoders) (0.17.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from scikit-learn>=0.20.0->category_encoders) (2.1.0)
Requirement already satisfied: geopandas in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (0.8.2)
Requirement already satisfied: fiona in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (1.8.18)
Requirement already satisfied: shapely in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (1.7.1)
Requirement already satisfied: pandas>=0.23.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (1.1.3)
Requirement already satisfied: pyproj>=2.2.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (3.0.0.post1)
Requirement already satisfied: cligj>=0.5 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona->geopandas) (0.7.1)
Requirement already satisfied: munch in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona->geopandas) (2.5.0)
Requirement already satisfied: attrs>=17 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona->geopandas) (20.2.0)
Requirement already satisfied: click-plugins>=1.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona->geopandas) (1.1.1)
Requirement already satisfied: click<8,>=4.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona->geopandas) (7.1.2)
Requirement already satisfied: six>=1.7 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona->geopandas) (1.15.0)
Requirement already satisfied: certifi in ./opt/anaconda3/envs/learn-env/
```

```
lib/python3.8/site-packages (from fiona->geopandas) (2020.6.20)
Requirement already satisfied: python-dateutil>=2.7.3 in ./opt/anaconda3/
envs/learn-env/lib/python3.8/site-packages (from pandas>=0.23.0->geopandas) (2.8.1)
Requirement already satisfied: pytz>=2017.2 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=0.23.0->geopandas) (2020.1)
Requirement already satisfied: numpy>=1.15.4 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=0.23.0->geopandas) (1.18.5)
```

```
In [310]: train['construction_year'].median()
```

```
Out[310]: 1986.0
```

```
In [311]: train['construction_year'] = train['construction_year'].replace({0:1986})
train['age'] = train['date_recorded'].astype(str).str[:4].astype(int) - train['pop/year'] = train['population'].replace({0:1}) / train['age'].repla
```

```
In [312]: X_test = test
```

```
In [313]: mean_year = full_df[full_df['construction_year']>0]['construction_year'].me
full_df.loc[full_df['construction_year']==0, 'construction_year'] = int(mea
X_test['age'] = X_test['date_recorded'].astype(str).str[:4].astype(int) - X
X_test['pop/year'] = X_test['population'].replace({0:1}) / X_test['age'].re
```

```
In [314]: train.head(10)
```

```
Out[314]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_id
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zai
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Ma
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahan Nany
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Si
5	9944	20.0	2011-03-13	Mkinga Distric Coun	0	DWE	39.172796	-4.765587	
6	19816	0.0	2012-10-01	Dwsp	0	DWSP	33.362410	-3.766365	Ng
7	54551	0.0	2012-10-09	Rwssp	0	DWE	32.620617	-4.226198	Tushir
8	53934	0.0	2012-11-03	Wateraid	0	Water Aid	32.711100	-5.146712	Ram
9	46144	0.0	2011-08-03	Isingiro Ho	0	Artisan	30.626991	-1.257051	Kw

10 rows × 42 columns

```
In [315]: train['water_/person'] = train['amount_tsh'].replace({0:1}) / train['popul
```

```
In [316]: X_test['water_/person'] = X_test['amount_tsh'].replace({0:1}) / X_test['po
```

```
In [318]: imputer = SimpleImputer(missing_values = np.nan, strategy = 'mean')
imputer.fit(features, labels['status_group'])
features = imputer.transform(features)
```

```
In [319]: scaler = RobustScaler()
scaler.fit(features, labels['status_group'])
features = scaler.transform(features)
```

```
In [320]: from pylab import rcParams
rcParams['figure.figsize'] = 30, 20
```

## OSEMN-Model

The model below is a rough scatterplot to showcase the different locations and functionalities of the wells geographically, it also pints a much more interesting picture of how certain locations are not having their needs met, which would certainly be something worth exploring in future work.

We will also be trying our hand at some logistic regression, and tree classification later on, but first le's fit our data to better serve our purpose.

```
In [321]: !pip install descartes
```

```
Requirement already satisfied: descartes in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (1.1.0)
Requirement already satisfied: matplotlib in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from descartes) (3.3.1)
Requirement already satisfied: kiwisolver>=1.0.1 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (1.2.0)
Requirement already satisfied: python-dateutil>=2.1 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (2.8.1)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.3 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (2.4.7)
Requirement already satisfied: pillow>=6.2.0 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (7.2.0)
Requirement already satisfied: numpy>=1.15 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (1.18.5)
Requirement already satisfied: cycler>=0.10 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (0.10.0)
Requirement already satisfied: certifi>=2020.06.20 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from matplotlib->descartes) (2020.6.20)
Requirement already satisfied: six>=1.5 in ./opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from python-dateutil>=2.1->matplotlib->descartes) (1.15.0)
```



```
In [322]: gdf = geopandas.GeoDataFrame(
            df, geometry=geopandas.points_from_xy(df.longitude, df.latitude))

functional = gdf.where(labels['status_group'] == 'functional')
repair = gdf.where(labels['status_group'] == 'functional needs repair')
broken = gdf.where(labels['status_group'] == 'non functional')

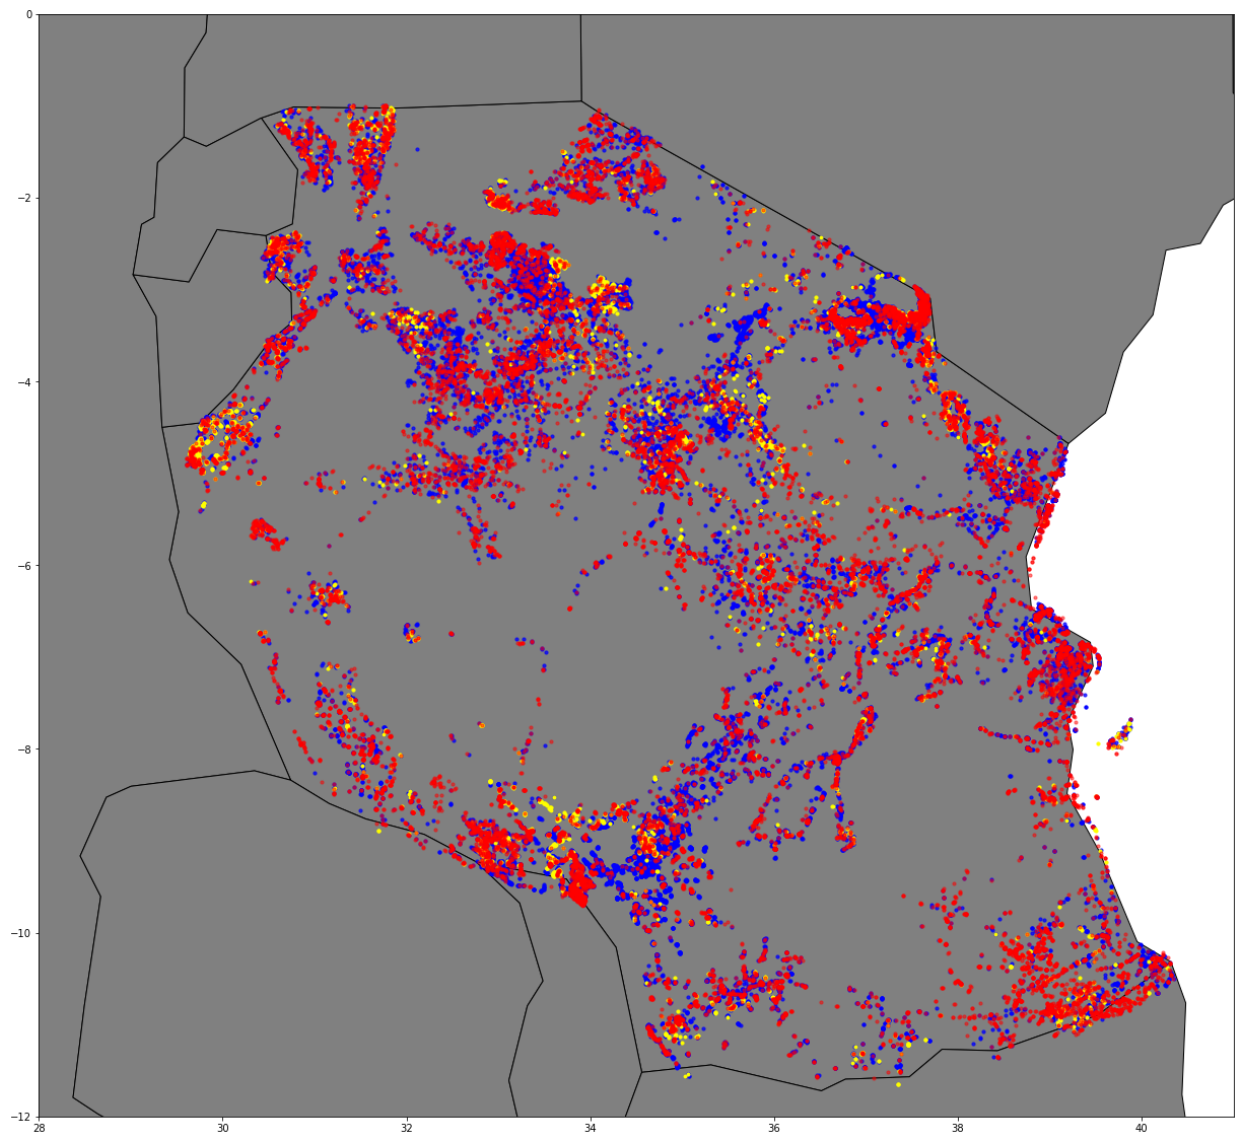
world = geopandas.read_file(geopandas.datasets.get_path('naturalearth_lowre

ax = world[world.continent == 'Africa'].plot(
    color='grey', edgecolor='black')
ax.scatter(functional['longitude'], functional['latitude'],
            c='blue', alpha=.75, s=10)
ax.scatter(repair['longitude'], repair['latitude'],
            c='yellow', alpha=1, s=10)
ax.scatter(broken['longitude'], broken['latitude'],
            c='red', alpha=.5, s=10)

plt.ylim(-12, 0)
plt.xlim(28, 41)

plt.show()

#blue is functional, yellow needs repair, red is broken
```



```
In [324]: full_df = pd.concat([train, test])  
full_df.shape  
#59400 + 14358
```

```
Out[324]: (74250, 43)
```

```
In [325]: full_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 74250 entries, 0 to 14849
Data columns (total 43 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     74250 non-null  int64
1   amount_tsh                           74250 non-null  float64
2   date_recorded                         74250 non-null  object
3   funder                                69746 non-null  object
4   gps_height                           74250 non-null  int64
5   installer                            69718 non-null  object
6   longitude                            74250 non-null  float64
7   latitude                             74250 non-null  float64
8   wpt_name                             74250 non-null  object
9   num_private                           74250 non-null  int64
10  basin                                74250 non-null  object
11  subvillage                           73780 non-null  object
12  region                               74250 non-null  object
13  region_code                          74250 non-null  int64
14  district_code                       74250 non-null  int64
15  lga                                  74250 non-null  object
16  ward                                 74250 non-null  object
17  population                           74250 non-null  int64
18  public_meeting                       70095 non-null  object
19  recorded_by                          74250 non-null  object
20  scheme_management                   69404 non-null  object
21  scheme_name                         38992 non-null  object
22  permit                             70457 non-null  object
23  construction_year                   74250 non-null  int64
24  extraction_type                     74250 non-null  object
25  extraction_type_group                74250 non-null  object
26  extraction_type_class                74250 non-null  object
27  management                           74250 non-null  object
28  management_group                    74250 non-null  object
29  payment                             74250 non-null  object
30  payment_type                        74250 non-null  object
31  water_quality                       74250 non-null  object
32  quality_group                       74250 non-null  object
33  quantity                            74250 non-null  object
34  quantity_group                      74250 non-null  object
35  source                              74250 non-null  object
36  source_type                         74250 non-null  object
37  source_class                        74250 non-null  object
38  waterpoint_type                     74250 non-null  object
39  waterpoint_type_group                74250 non-null  object
40  age                                 74250 non-null  int64
41  pop/year                            74250 non-null  float64
42  water_/person                       74250 non-null  float64
dtypes: float64(5), int64(8), object(30)
memory usage: 24.9+ MB
```

```
In [326]: for column in full_df.columns:
          full_df[column].fillna(full_df[column].mode()[0], inplace=True)
```

```
In [327]: full_df.isna().sum()
```

```
Out[327]: id                                0
          amount_tsh                        0
          date_recorded                    0
          funder                            0
          gps_height                        0
          installer                         0
          longitude                         0
          latitude                          0
          wpt_name                          0
          num_private                       0
          basin                             0
          subvillage                       0
          region                           0
          region_code                      0
          district_code                    0
          lga                              0
          ward                             0
          population                       0
          public_meeting                   0
          recorded_by                      0
          scheme_management                 0
          scheme_name                      0
          permit                           0
          construction_year                0
          extraction_type                  0
          extraction_type_group            0
          extraction_type_class            0
          management                       0
          management_group                 0
          payment                          0
          payment_type                     0
          water_quality                    0
          quality_group                    0
          quantity                         0
          quantity_group                   0
          source                           0
          source_type                      0
          source_class                     0
          waterpoint_type                  0
          waterpoint_type_group            0
          age                              0
          pop/year                         0
          water_/person                    0
          dtype: int64
```

```
In [254]: test = (test.drop(['date_recorded'], axis = 1),
                  test.drop(['construction_year'], axis = 1))
```

```
In [256]: labels.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   id              59400 non-null  int64
 1   status_group    59400 non-null  object
dtypes: int64(1), object(1)
memory usage: 928.2+ KB
```

```
In [259]: X_matrix, y_vector = train, labels['status_group']
```

```
In [260]: majority_class = y_vector.mode()
y_vector.value_counts(normalize=True)
```

```
Out[260]: functional          0.543081
non functional              0.384242
functional needs repair     0.072677
Name: status_group, dtype: float64
```

```
In [261]: majority_prediction = [majority_class] * len(y_vector)
accuracy_score(y_vector, majority_prediction)
```

```
Out[261]: 0.543080808080808
```

```
In [265]: X_cleaned = full_df[:-14358]
X_test_cleaned = full_df[-14358:]
y = labels['status_group']
```

```
In [266]: X_cleaned.shape, X_test_cleaned.shape, y.shape
```

```
Out[266]: ((59892, 43), (14358, 43), (59400,))
```

```
In [270]: X_matrix.head()
y_vector.head()
```

```
Out[270]: 0      functional
1      functional
2      functional
3  non functional
4      functional
Name: status_group, dtype: object
```

```
In [271]: X_matrix.dtypes
```

```
Out[271]: id                int64
amount_tsh                float64
date_recorded             object
funder                    object
gps_height                int64
installer                 object
longitude                 float64
latitude                  float64
wpt_name                  object
num_private               int64
basin                     object
subvillage                 object
region                    object
region_code               int64
district_code             int64
lga                       object
ward                      object
population                int64
public_meeting            object
recorded_by               object
scheme_management         object
scheme_name               object
permit                    object
construction_year         int64
extraction_type            object
extraction_type_group     object
extraction_type_class     object
management                object
management_group          object
payment                   object
payment_type              object
water_quality             object
quality_group             object
quantity                  object
quantity_group            object
source                    object
source_type               object
source_class              object
waterpoint_type           object
waterpoint_type_group     object
age                       int64
pop/year                  float64
water_/_person            float64
dtype: object
```

```
In [377]: len(y_vector)
```

```
Out[377]: 14850
```

```
In [275]: X_train_numeric = X_matrix.select_dtypes(np.number)
X_test_numeric = X_test.select_dtypes(np.number)
```

```
In [281]: train.head()
```

```
Out[281]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	r
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zah
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Mah
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zah
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Nanyu

5 rows × 43 columns

```
In [328]: majority_class = labels['status_group'].mode()[0]
y_pred = np.full(shape = labels['status_group'].shape, fill_value = majority_class)
```

```
In [329]: accuracy_score(labels['status_group'], y_pred)
```

```
Out[329]: 0.543080808080808
```

```
In [332]: print(classification_report(labels['status_group'], y_pred))
#oof that could look a lot better...
```

	precision	recall	f1-score	support
functional	0.54	1.00	0.70	32259
functional needs repair	0.00	0.00	0.00	4317
non functional	0.00	0.00	0.00	22824
accuracy			0.54	59400
macro avg	0.18	0.33	0.23	59400
weighted avg	0.29	0.54	0.38	59400

```
In [369]: len(X_train_numeric)
```

```
Out[369]: 59400
```

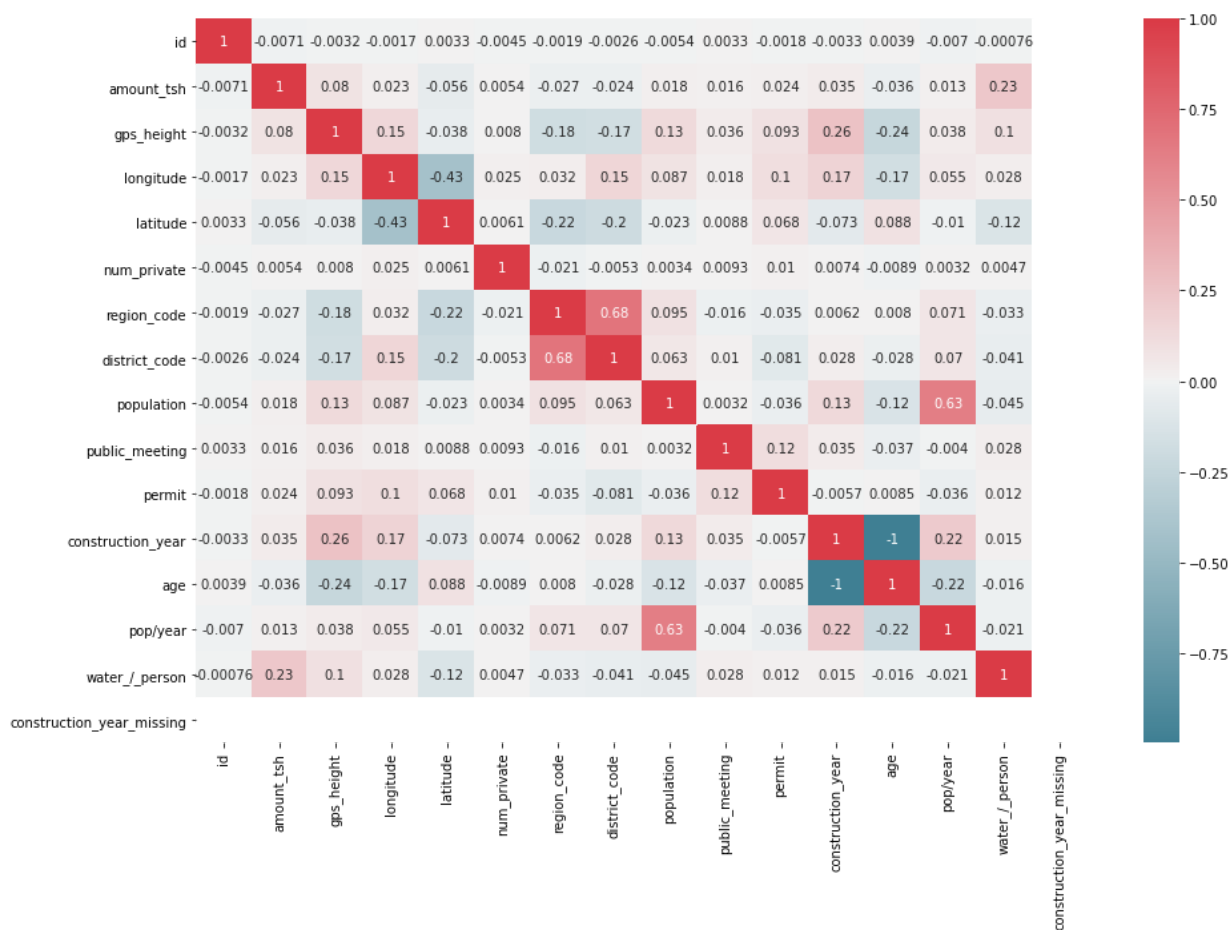
```
In [370]: X_matrix, X_train_numeric, y, y_vector = train_test_split(X_train_numeric,
```

```
In [398]: clf = tree.DecisionTreeClassifier()
clf.fit(X = X_matrix, y = y)
clf.feature_importances_
clf.score(X = X_test_numeric, y = y_vector)
```

Out[398]: 0.44195286195286193

```
In [423]: def correlation_heatmap(labels):
ax=plt.subplots(figsize=(15,10))
colormap=sns.diverging_palette(220,10,as_cmap=True)
sns.heatmap(full_df.corr(),annot=True,cmap=colormap)
```

correlation\_heatmap(labels)



```
In [426]: #new_df = pd.concat([train, labels])
new_df = pd.merge(train, labels, on = 'id')
```



```
In [421]: new_df.head()
```

```
Out[421]:
```

	num_private	...	quantity_group	source	source_type	source_class	waterpoint_type	waterpoint
1e	0	...	enough	spring	spring	groundwater	communal standpipe	commu
ati	0	...	insufficient	rainwater harvesting	rainwater harvesting	surface	communal standpipe	commu
/a di	0	...	enough	dam	dam	surface	communal standpipe multiple	commu
ati /a ou	0	...	dry	machine dbh	borehole	groundwater	communal standpipe multiple	commu
ni	0	...	seasonal	rainwater harvesting	rainwater harvesting	surface	communal standpipe	commu

```
In [349]: y.shape, X_matrix.shape
```

```
Out[349]: ((59400,), (59400, 42))
```

```
In [396]: len(y_vector)
```

```
Out[396]: 14850
```

## OSEMN-Interpret

Our model didn't exactly live up to our expectations, 44% would be worse than flipping a coin, especially considering that over 50% of the wells in Tanzania are currently operational. We did however get to take a look at geographic distributions, and took an opportunity to experiment with logistic regression and classification.

Unfortunately, the interpretation of this data would be more speculative than most any analysts would be comfortable with. The interpretation of this data could be most accurately described as needing further work, if we want to achieve actionable results. However, this will provide opportunities to practice ternary classification and predictive modeling, which can help solve future problems.

```
In [ ]:
```

